

# **3D Soaring through unknown Atmosphere**

**Benjamin Rothaupt**

**Masterarbeit**

**2018**



Universität Stuttgart  
Institut für Flugmechanik und Flugregelung

---



# Erklärung

Hiermit versichere ich, dass ich diese Masterarbeit selbständig mit Unterstützung der Betreuer angefertigt und keine anderen als die angegebenen Quellen und Hilfsmittel verwendet habe.

Die Arbeit oder wesentliche Bestandteile davon sind weder an dieser noch an einer anderen Bildungseinrichtung bereits zur Erlangung eines Abschlusses eingereicht worden.

Ich erkläre weiterhin, bei der Erstellung der Arbeit die einschlägigen Bestimmungen zum Urheberschutz fremder Beiträge entsprechend den Regeln guter wissenschaftlicher Praxis<sup>1</sup> eingehalten zu haben. Soweit meine Arbeit fremde Beiträge (z.B. in der Form von Bildern, Zeichnungen, Textpassagen etc.) enthält, habe ich diese Beiträge als solche gekennzeichnet (Zitat, Quellenangabe) und eventuell erforderlich gewordene Zustimmungen der Urheber zur Nutzung dieser Beiträge in meiner Arbeit eingeholt. Mir ist bekannt, dass ich im Fall einer schuldhaften Verletzung dieser Pflichten die daraus entstehenden Konsequenzen zu tragen habe.

---

Ort, Datum

---

Benjamin Rothaupt

**Betreuer: M. Sc. Stefan Notter**

---

<sup>1</sup>Nachzulesen in den DFG-Empfehlungen zur „Sicherung guter wissenschaftlicher Praxis“ bzw. in der Satzung der Universität Stuttgart zur „Sicherung der Integrität wissenschaftlicher Praxis und zum Umgang mit Fehlverhalten in der Wissenschaft“



## **Kurzfassung**

Testtext Testtext 123

## **Abstract**

Dies ist der englische Abstract.



# Contents

<b>Erklärung</b>	<b>iii</b>
<b>Kurzfassung / Abstract</b>	<b>v</b>
<b>List of Figures</b>	<b>ix</b>
<b>List of Tables</b>	<b>xi</b>
<b>Nomenklatur</b>	<b>xiii</b>
<b>1. Introduction</b>	<b>1</b>
<b>2. Reinforcement Learning</b>	<b>3</b>
2.1. The Agent - Environment System . . . . .	3
2.2. Model Based and Model Free Learning . . . . .	3
2.3. Agent . . . . .	4
2.4. Return . . . . .	5
2.5. Markov Decision Process . . . . .	5
2.6. Policy . . . . .	6
2.7. Value Functions . . . . .	6
2.8. The Bellman Expectation Equation . . . . .	7
2.9. The Bellman Optimality Equation . . . . .	8
<b>3. Function Approximation</b>	<b>9</b>
3.1. Tables . . . . .	9
3.2. Artificial Neural Networks . . . . .	10
3.3. Supervised Learning . . . . .	14
3.4. Optimization Techniques . . . . .	14
3.4.1. Gradient Descent . . . . .	14
3.4.2. Stochastic Gradient Descent . . . . .	15
3.4.3. Momentum . . . . .	15
3.4.4. The ADAM Algorithm . . . . .	16
3.5. Overfitting . . . . .	16
<b>4. Dynamic Programming</b>	<b>19</b>
4.1. The Principle of Optimality . . . . .	21

## Contents

4.2. Types of Dynamic Programming Algorithms . . . . .	21
4.2.1. Policy Evaluation . . . . .	21
4.2.2. Policy Iteration . . . . .	22
4.2.3. Value Iteration . . . . .	23
4.3. The Contraction Mapping Theorem . . . . .	23
<b>5. Trajectory Optimization with Policy Iteration</b>	<b>25</b>
5.1. Glider Representation . . . . .	25
5.2. 2D Environment . . . . .	25
5.2.1. Policy Representation . . . . .	25
5.2.2. Equations of Motion . . . . .	26
5.2.3. Maximum Range in a given Configuration . . . . .	27
5.2.4. Following the Line of Sight Towards the Target . . . . .	28
5.2.5. Discretization of the State- and Action Space . . . . .	28
5.3. 3D Environment . . . . .	28
5.3.1. Discretization of the state and action space . . . . .	28
<b>6. Results</b>	<b>31</b>
6.1. 2D Policy Iteration . . . . .	31
6.2. 3D Scenarios . . . . .	33
<b>7. Discussion</b>	<b>35</b>
<b>A. Appendix A: Glider Parameters and Scenario Data</b>	<b>37</b>
<b>B. Appendix B: Computer Configuration</b>	<b>39</b>
<b>C. Appendix C: Standardization Values</b>	<b>41</b>
C.1. 2D Scenarios . . . . .	41
<b>Literaturverzeichnis</b>	<b>43</b>



# List of Figures

2.1. The Agent-Environment-System [6] . . . . .	4
3.1. A neuron from an artificial neural network [6] . . . . .	11
3.2. The MLP used for the policy [6] . . . . .	12
4.1. classification of Machine Learning algorithms . . . . .	19
4.2. Figure 3.1 von Markus . . . . .	20
5.1. Angles in three-dimensional flight. . . . .	26
5.2. The discretized state space. For one set of $x_3$ and $z_4$ , the grid for the speed vector is drawn. . . . .	29
6.1. Trajectory and angle of attack after 100 Optimistic Policy Iterations . . .	31
6.2. Trajectory and angle of attack after 100 Generalized Policy Iterations with up to 100 evaluation steps in each iteration . . . . .	32
6.3. Trajectory and angle of attack after 100 Value Iterations . . . . .	33
6.4. Trajectory and angle of attack after 100 iterations of Generalized Policy Iteration with 100 evaluation steps in each iteration . . . . .	34



# List of Tables

5.1. Grid parameters for trajectory optimization . . . . .	29
A.1. glider data . . . . .	37
A.2. scenario data . . . . .	38
B.1. computer specs . . . . .	39
C.1. Standardization values for 2D trajectory optimization . . . . .	41



# Nomenclature

**Abbreviations**

**Latin Letters**

**Greek Letters**

**Indices**



# 1. Introduction

In autonomous soaring, it is vital to gather as much energy from the atmosphere as possible in order to stay aloft for long time periods. Generally, the locations where updrafts occur are not known a-priori. So a pilot has to react spontaneously to changes in his vicinity in order to harvest as much energy as possible, which takes a lot of experience.

Electric aircraft have been an important research topic in recent years. They can operate environmentally friendly and are easy to maintain. Their range and endurance is - however - limited by the energy density of current batteries. This means they either have to land frequently to recharge or be equipped with solar panels on their wings which weigh them down and are difficult to maintain. Electric UAVs suffer from the same deficiency. They can theoretically be used for a wide range of applications from gathering scientific data to surveillance, battery capacity being their only limiting factor. The exploitation of thermal updrafts is one way to increase the endurance of an aircraft without changing anything on the airframe itself which makes it relatively cheap to implement.

Finding an optimal flight path with respect to varying mission goals in unknown atmosphere is a challenging task. While the dynamics of the aircraft itself are known and understood, the occurrence of updrafts cannot be predicted with sufficient accuracy. The task of finding an optimal path in moving air can be split into finding an optimal path in calm air and dealing with the unknown updrafts separately. Reinforcement Learning algorithms are very good at finding an optimal path through an unknown environment. They require no prior knowledge and learn from experience they gather through simulations. This makes them well suited for learning how to react optimally to random updrafts.

Finding an optimal path through calm air can be done more efficiently, utilizing the knowledge about how a glider behaves in calm air. Flying a glider can be regarded as a time-discrete decision process where, at each time step  $t$ , the pilot selects a control surface position. The states  $s_t$  in this decision process have the Markov property, i.e. the information needed to calculate  $s_{t+1}$  is fully contained in  $s_t$ . A decision process where the Markov property applies to all states is called a Markov Decision Process (MDP). Dynamic Programming is one of the most popular machine learning algorithms for MDPs with up to a few millions of states.

Dynamic Programming is a general term that applies to problems that can be split into subproblems. The optimal solution of each subproblem can then be used to re-compose the optimal solution to the whole problem. If the subproblems are similar, DP is especially effective. DP however suffers from the "curse of dimensionality", in large or continuous state spaces it quickly becomes inefficient. The obvious approach is to dis-

## *1. Introduction*

cretize the state space and deal with a smaller number of states. If done properly, this yields an approximation of the exact optimal solution. Depending on the problem and the discretization, the achieved approximate solution is more or less close to the exact one. Like in many technical topics, there is a tradeoff between precision and cost.

In this work, a time optimal flight path through calm air is calculated with dynamic programming. First, the state and action space are discretized. A policy iteration algorithm is then used to calculate optimal paths in the vertical plane and in a 3D environment. The usefulness of the results from the 3D scenario is obvious. The results from the 2D scenario are also useful where the ground trace of a trajectory is dictated by boundary conditions such as topography or other airspace restrictions.



## 2. Reinforcement Learning

### 2.1. The Agent - Environment System

In Reinforcement Learning scenarios, an agent interacts with its environment and thereby tries to maximize some sort of reward. By interacting with its environment, it learns how to behave optimally with respect to a scalar reward function. The agent aims to maximize its total reward by choosing - at each time step - the best possible action  $a$  from a set of actions  $\mathcal{A}$ . The agent has no prior knowledge about the environment and about how to behave in an optimal way. Instead, it learns by interacting with its environment and the information it receives during training. Before each step, the agent is in a state  $s_t$  in the environment, chooses an action  $a_t$  and receives a scalar reward  $r_t$  and an observation  $o_t$  belonging to the next state  $s_{t+1}$  and the process is repeated. Figure 2.1 illustrates what happens at each time step. By interacting with the environment, the agent learns which action is optimal for each state it visits.

The value of  $r_t$  the agent receives at each step is determined by a so called reward function. The choice of a reward function is crucial for the success of any RL-based trajectory optimization algorithm. The only way to "tell" the agent, what his goal is, is by the reward function. Generally, a reward function should give the agent a positive reward for actions that bring him closer to his goal and penalize actions that reduce his chances of reaching it. If an action takes the agent to a terminal state that is not the goal, this action should yield a big penalty. Penalties are a means of keeping the agent away from states that are not desirable, such as unwanted terminal states or states in their direct vicinity.

### 2.2. Model Based and Model Free Learning

MDP solution methods can be divided into model-free and model-based methods. In model-based methods, the agent uses a model to predict the reaction of the environment to any of his actions. An environment model typically takes a state-action pair and returns the next state and next reward. If the environment is stochastic, there are multiple possible next states and rewards. A *sample model* returns one of the possibilities, whereas a *distribution model* returns some representation of all possible next states and rewards.

Regardless of the type, all models simulate what might happen to the agent in future time steps. Whether the experience gathered through a model is useful obviously depends on its quality. To distinguish experience from a model from real experience, the results

## 2. Reinforcement Learning

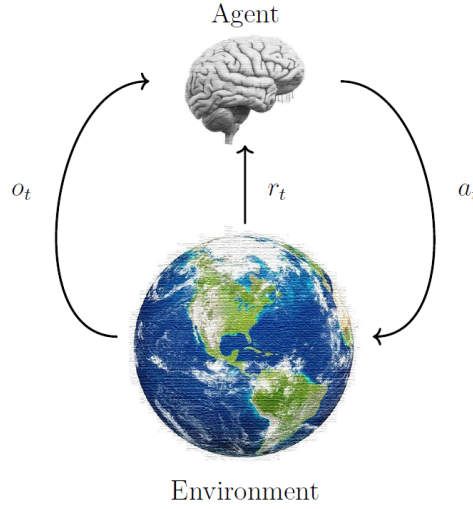


Figure 2.1.: The Agent-Environment-System [6]

from a model are referred to as *simulated experience*[5]. Figure 4.2 shows some popular Machine Learning algorithms and classifies them.

### 2.3. Agent

The agent in an MDP consists of at least one of the following parts.

- Policy:  
A mapping from states to actions that tells the agent what to do (see section 2.6)
- Value Function:  
A mapping from states (or state-action pairs) to values (see section 2.7)
- Model:  
A model of the environment that might be updated by the agent's experience

Most agents have a policy. They also can have a value function, like in actor-critic algorithms where the value function is used for bootstrapping. This means utilizing the Bellman Expectation Equation to estimate the return  $G_t$  by doing a one step lookahead and use the sum of the immediate reward and the successor state value.

$$\mathbb{E}[G_t | s_t = s] = \mathbb{E}[R_{t+1} + \gamma G_{t+1} | s_t = s, a_t = a] \quad (2.1)$$

$$= \mathbb{E}[R_{t+1} + \gamma V(s_{t+1}) | s_t = s, a_t = a] \quad (2.2)$$

Instead of sampling a complete episode, bootstrapping makes it possible to learn from incomplete episodes.

A model of the environment is - in principle - not required for successful training. It can however be used to gather simulated experience and therefore reduce training time on a real agent like a robot. Although imperfect, information from a model can be useful in such cases where training time is restricted.

## 2.4. Return

The return  $G_t$  at a time step  $t$  is the cumulative reward  $r_t$  at each step from the state  $s_t$  onwards until reaching a terminal state:

$$G_t = \sum_{k=0}^{T-t-1} \gamma^k r_{t+k+1} \quad (2.3)$$

If  $\gamma \in (0, 1]$  is close to zero, immediate rewards are weighted more. If  $\gamma$  is close to one, rewards in the distant future are weighted more, making decisions more far sighted. In infinite horizon problems where there is no terminal state,  $\gamma$  must be less than one in order to avoid infinite returns. In finite horizon (i.e. episodic) MDPs,  $\gamma$  is usually close to or exactly one.

## 2.5. Markov Decision Process

A Markov Decision Process is a series of decision problems where each decision has an impact on the final reward the agent receives. In an MDP, the probability of reaching a state  $s_{t+1}$  only depends on the state  $s_t$ . The predecessors of  $s_t$  are irrelevant for  $P(s_{t+1}|s_t, a_t)$ .

$$P(s', r|s, a) = \mathbb{P}[s_{t+1} = s', r_{t+1} = r | s_t = s, a_t = a] \quad (2.4)$$

$$= \mathbb{P}[s_{t+1} | s_t, a_t, s_{t-1}, a_{t-1}, s_{t-2}, a_{t-2}, \dots] \quad (2.5)$$

Every decision process where this probability does not only depend on the direct predecessor can be made markovian by defining additional state variables representing the states that came before. This extended current state contains all relevant information for the transition to  $s_{t+1}$ .

In most RL-problems, the goal is to maximize the total return from a given initial state. This is done by trying to find the optimal policy  $\pi^*$  which yields the maximum expected return from each state.  $\pi^*$  can be achieved directly or by finding the optimal value function, i.e. the mapping from each state or state-action-pair to its true value  $v^*$ . If  $\pi^*$  is found for a specific MDP, the MDP is solved.

## 2.6. Policy

In reinforcement learning problems, a policy is the central part of the agent. It contains all the information needed to decide what action to choose in each state the agent visits. Mathematically, a policy is a mapping from states to actions.

$$\pi : \mathcal{S} \times \mathcal{A} \mapsto [0, 1] \quad (2.6)$$

The policy is a probability distribution over actions given states and returns - at each state - the probability of taking an action  $a_t$ . If the policy is deterministic, only one  $a_t$  out of  $\mathcal{A}$  has a probability of one. In a stochastic policy, the sum of all possible action-probabilities must be one.

$$\pi(a|s) = \mathbb{P}[a_t = a | s_t = s] \quad (2.7)$$

This yields equations 2.8 and 2.9 for the state transition probability and the reward function under policy  $\pi$ .

$$\mathcal{P}^\pi(s, s') = \sum_a \pi(a|s) \mathcal{P}(s'|s, a) \quad (2.8)$$

$$\mathcal{R}^\pi(s) = \sum_a \pi(a|s) \mathcal{R}(s, a) \quad (2.9)$$

At each timestep, the agent passes its state observation to the policy and the policy returns an action. This action is taken and the agent gets a new observation. This process is repeated until the agent reaches a terminal state.

A policy can be implemented in many different ways. The simplest option is a table where each field contains the action for one state or - if the policy is stochastic - all possible actions and their respective probabilities. Other examples include linear combinations of the inputs and linear combinations of basis functions. The most popular way to implement a policy is - however - through an artificial neural network (ANN). More information about neural networks is given in chapter 3.

## 2.7. Value Functions

Some RL-algorithms try to find a policy directly from experience. Others, like A3C, additionally utilize a so called value function. A value function maps from states or actions to their values. In this context, the value of a state or an action is a measure of how useful it is for the agent to visit the respective state or choosing the respective action. There are two types of value functions, state value functions and action value functions.

$$V(s) = \mathbb{E}[G_t | s_t = s] \quad (2.10)$$

## 2.8. The Bellman Expectation Equation

$$Q(a|s) = \mathbb{E}[G_t | s_t = s, a_t = a] \quad (2.11)$$

In the following sections, the action space is regarded finite for simplicity. The state value function 2.10 can be expressed in terms of the action value function 2.11:

$$V_\pi(s) = \sum_a \pi(a|s) Q_\pi(s, a) \quad (2.12)$$

Similarly, the action value function can be expressed by means of the state value function:

$$Q_\pi(s, a) = \mathcal{R}(s, a) + \gamma \sum_{s'} \mathcal{P}(s'|s, a) V_\pi(s') \quad (2.13)$$

If an agent has found the optimal value function  $V^*$  or  $Q^*$  of an MDP, the optimal policy  $\pi^*$  can be derived directly by acting greedily with respect to the value function at each state.

Like a policy, a value function can be implemented with a table, linear combinations of the inputs, basis functions and an artificial neural network.

## 2.8. The Bellman Expectation Equation

According to the Bellman Expectation Equation, every path through the state space can be decomposed into two parts, the next step and the rest of the path. Equivalently, every state value can be split into the value of the next step and the successor state value.

The state value function 2.10 can be written as a weighted sum of the action values each multiplied with the probability to take the respective action according to the policy. Replacing  $Q_\pi(s, a)$  in equation 2.12 with equation 2.13 yields equation 2.14 which is known as the *bellman expectation equation*:

$$V_\pi(s) = \sum_a \pi(a|s) Q_\pi(s, a) = \sum_a \pi(a|s) \sum_r \sum_{s'} \mathcal{P}(s', r|s, a) [r + \gamma V_\pi(s')] \quad (2.14)$$

In a deterministic environment, equation 2.14 becomes

$$V_\pi(s) = \sum_a \pi(a|s) [r + \gamma V_\pi(s')] \quad (2.15)$$

The Bellman Expectation Equation can be used to update a state value by looking at the state values of its successor states and weighting them by the probabilities to reach each successor state (see equations 4.6 and 4.7).

## 2.9. The Bellman Optimality Equation

Instead of a weighted sum of rewards and successor values, like in equation 2.14, one can take only the action that has the maximum action value. This yields the following equation:

$$v_*(s) = \max_{a \in \mathcal{A}(s)} q_{\pi_*}(s, a) \quad (2.16)$$

$$= \max_a \mathbb{E}_{\pi_*}[G_t | s_t = s, a_t = a] \quad (2.17)$$

$$= \max_a \mathbb{E}_{\pi_*}[R_{t+1} + \gamma G_{t+1} | s_t = s, a_t = a] \quad (2.18)$$

$$= \max_a \mathbb{E}[R_{t+1} + \gamma V_*(s_{t+1}) | s_t = s, a_t = a] \quad (2.19)$$

$$= \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V_*(s')] \quad (2.20)$$

If the environment is deterministic, equation 2.20 can be simplified:

$$v_*(s) = \max_a [r + \gamma v_*(s')] \quad (2.21)$$

All equations from 2.16 to 2.21 are alternative ways to express the *bellman optimality equation*. It states that the value of a state under an optimal policy must equal the expected return for the best action in that state [5]. Intuitively it is obvious that the optimal policy must pick the best action in each state. In Value Iteration, the *bellman optimality equation* is used as an update rule. See chapter 4.2.3 for more details.

## 3. Function Approximation

A common problem in all engineering topics is developing a model from measurement data. Such a model makes it possible to generalize from the given data to new data where a real measurement does not exist. The process of developing and fitting it to given data is called function approximation. It is essential for topics like image recognition, but also regression.

Mathematically, function approximation is the process finding a model function that maps from input variables  $x$  to outputs  $y$ . Usually, a set of  $x_{fit}$  and corresponding  $y_{fit}$  is given and the goal is to find a function that gives a value close to  $y_{fit,i}$  for each  $x_{fit,i}$ .

In the context of machine learning, artificial neural networks are commonly used for classification, clustering and regression. In classification, the neural network is trained on a set of samples, each belonging to a category which is supplied to the network. As there typically is a finite number of categories, the output of the network is discrete. The training aims to enable the network to put samples it has not been trained on in the correct category. In regression, the network is trained to approximate a function that maps from the input(-vector) to a continuous output(-vector). In this work, for example, the policy is represented as an artificial neural network. In clustering, the network is given a set of training samples without explicitly knowing what cluster they belong to. The training goal is for the network to find appropriate clusters and put the samples into the right cluster so that similar samples are in the same cluster.

### 3.1. Tables

The simplest way to express a discrete mapping is by a table where each input/output pair is represented by a row of the table. The number of inputs and corresponding outputs in a table is inherently finite. The set of states and actions in a continuous MDP is however infinite, so the policy and value functions must also be continuous. If a discrete mapping like a table is used as a representation of a policy or value function in a continuous MDP, it will, at some point, be confronted with unknown inputs. The inability to deal with unknown inputs is a big disadvantage of tables. There are, however, ways to overcome this weakness.

(Hier noch Tabelle für Value Function und Policy bzw. Greedy Actions einfügen)  
(Erklären, dass die vordere Spalte, wenn sie zustände enthält, auch aus mehreren Spalten bestehen kann, was beim berechnen der nearest neighbor oder linearer interpolation erforderlich ist)

### 3. Function Approximation

#### Nearest Neighbor

Let  $f : \mathbb{R} \mapsto \mathbb{R}, y = f(x)$  be a continuous mapping from a scalar to a scalar.

The most straightforward way to obtain an output for an unknown input is to look for the table entry that is closest to the unknown input. This yields a piecewise constant mapping from inputs to outputs and enables the table to generalize. Whether this suffices as an approximation for a continuous mapping  $f$  heavily depends on the properties of  $f$ . If there are large differences between the output of two subsequent inputs, the error can be quite high.

#### Linear Interpolation

Another way to obtain outputs for unknown inputs is to interpolate linearly. In a first step, the two nearest known inputs  $x_0$  and  $x_1$  and their respective outputs  $y_0$  and  $y_1$  are looked up (note that  $x_0 \leq x \leq x_1$  in this case). In the second step, the output is calculated by the following formula:

$$y_{interp} = y_0 + \frac{y_1 - y_0}{x_1 - x_0} \cdot (x - x_0) \quad (3.1)$$

One advantage of linear interpolation of table values is that the output is at least continuous, though not differentiable at the known input points.

An alternative formulation of equation 3.1 is:

$$y_{interp} = y_0 \cdot (1 - \gamma) + y_1 \cdot \gamma \quad (3.2)$$

where  $\gamma = \frac{x - x_0}{x_1 - x_0}$ .

## 3.2. Artificial Neural Networks

The idea behind artificial neural networks is to mimic the structure of a human brain. The human brain consists of a very high number of neurons, each receiving signals from other neurons, processing them and itself sending a signal to other neurons if certain conditions are met. Each neuron is very simple, it is the high number of them and the fact that they all work in parallel that makes the human brain so powerful.

A neuron in an artificial neural network looks like 3.1. It takes an arbitrary number of inputs, multiplies each input by the corresponding input weight, adds a bias and applies to the result a nonlinear function, the activation function. Multiple activation functions are common, but they all are somewhat s-shaped. Examples include the step function, ( $f(x) = 1$  if  $x \geq x_0$ ,  $f(x) = 0$  otherwise), the rectifier nonlinearity ( $f(x) = \max(0, x)$ ), the sigmoid function ( $f(x) = 1/(1 + e^{-x})$ ) and the tanh-activation function ( $f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$ ).

In an artificial neural network, the neurons are arranged in layers. The ANNs in this work are all fully connected feed forward networks, also known as Multi Layer Perceptrons






Figure 3.1.: A neuron from an artificial neural network [6]

(MLPs). Each neuron in a given layer  $k$  gets an input from each neuron in the previous layer  $k - 1$  and outputs a signal to each neuron in the next layer  $k + 1$  (see figure 3.1).

Modern computer GPUs also have a high number of processing units and are therefore well suited for parallelized calculations.

If a policy or value function is represented by an artificial neural network (ANN), the output for each state (or state-action-pair) is determined by the weights and biases of the network. One set of weights and biases forms the *parameter vector*  $\theta$  of the neural network. If, for example, a stochastic policy is parametrized by  $\theta$  and represented by a neural net with  $k$  layers, the notation is as follows:

$$\pi_{\theta}(a|s) = \mathbb{P}[a|s, \theta] \quad (3.3)$$

with

$$\theta = \{\mathbf{W}^{(1)}, \mathbf{b}^{(1)}, \mathbf{W}^{(2)}, \mathbf{b}^{(2)}, \dots, \mathbf{W}^{(k)}, \mathbf{b}^{(k)}\} \quad (3.4)$$

### 3. Function Approximation




Figure 3.2.: The MLP used for the policy [6]

The policy ANN does not necessarily need to output the action directly. Instead, most RL-related policies are stochastic, i.e. the neural net outputs the mean and standard deviation of a normal distribution:

$$\pi_{\theta}(\mu_{\theta}, \sigma_{\theta} | s) = \frac{1}{\sqrt{2\pi}\sigma_{\theta}} e^{-\frac{(a-\mu_{\theta})^2}{2\sigma_{\theta}^2}} \quad (3.5)$$

The action  $a$  is then drawn from  $\mathcal{N}(\mu_{\theta}(s), \sigma_{\theta}^2(s))$ .

#### Input Standardization

All training algorithms that are used in this work rely on the partial derivative of a loss function with respect to the network weights and biases. As can be seen in 3.8 and the following equations, they perform a step in (an approximation of) the current gradient direction. This gradient heavily depends on the network inputs. If some of the inputs have much larger values than the rest, this can result in some weights being updated

faster than others, which is not desirable. Therefore, the inputs to a neural network should have the same scale. In this work, the inputs are scaled to a normal distribution with  $\mu = 0$  and  $\sigma = 1$ :

$$\hat{x} = \frac{x - \mu}{\sigma} \quad (3.6)$$

The specific values for  $\mu$  and  $\sigma$  for each of the scenarios are presented in table C.1.

## Action Scaling

In an artificial neural network, every neuron can output a value between -1 and 1. Therefore, the output of a neural network is limited by the number of neurons in the last hidden layer. Obtaining a large output requires a lot of neurons. On the other hand, lots of neurons favor overfitting and increase training time. If the desired outputs are too small, then so are the neuron outputs. The tanh activation function that is used for the policy is approximately linear in this area. Therefore the whole network behaves almost linearly if it is trained to output values close to zero and linear behavior can also be obtained without any hidden layers. The most important advantage of deep neural networks is however their ability to represent highly nonlinear behavior.

Both problems can be avoided if outputs lie between -1 and 1. This way it is unlikely that any of the activation functions reach their saturation limit, nor is the network restricted to linear behavior. Therefore, the neural network is usually trained to output values between -1 and 1 and these values are then scaled to the desired magnitude by multiplying them with a scaling factor.

## Weight initialization

Training an artificial neural network is an iterative process. Therefore it requires initial values for the weights and biases. As training a neural network through supervised learning is - in principle - equivalent to solving an optimization problem iteratively, training success highly depends on the initial weights.

The two most popular ways to initialize neural network weights are the glorot uniform initialization and the glorot normal distribution. According to the glorot uniform initialization, the weights in one layer must be distributed uniformly within  $\pm \sqrt{\frac{6}{n_{in} + n_{out}}}$  where  $n_{in}$  and  $n_{out}$  are the numbers of neurons in the prior and current layer, respectively.

The glorot normal initialization states that the weights must be initialized according to a normal distribution with zero mean and a variance of  $\sigma = \sqrt{\frac{2}{n_{in} + n_{out}}}$ .

In this work, all neural networks are initialized according to the glorot uniform distribution.

## 3.3. Supervised Learning

All machine learning algorithms can be divided into two categories, supervised learning and unsupervised learning. In supervised learning, the agent is given the desired action for each state explicitly, so it knows what it is supposed to do. In unsupervised learning, however, the agent does not know what the best action is so it has to find out itself. One way of doing that is from moving in the state space and learning from experience, for example by applying a reinforcement learning algorithm.

In the supervised learning case, training is done via minimizing a loss function that depends on the desired action and the actual action taken by the agent. A common loss-function is the mean squared error (MSE):

$$L = \frac{1}{N} \cdot \sum_{n=1}^N L_n = \frac{1}{N} \cdot \sum_{n=1}^N \left[ \frac{1}{2} \cdot (a_n(s) - a_{n,desired})^2 \right] \quad (3.7)$$

The goal of supervised learning is to minimize the difference between  $a_{n,desired}$  and  $a_n(s)$  in each state. If  $a_n(s)$  is the output of a function approximator, minimizing  $L_{MSE}$  is equivalent to fitting the function approximator to the desired outputs. If  $L_{MSE}$  is minimal,  $a_n(s)$  is as close as possible to  $a_{n,desired}$  in most states.

One advantage of the application of a loss function is that numerous iterative optimization algorithms have been developed to solve such a problem.

## 3.4. Optimization Techniques

As mentioned before, supervised learning is usually done by defining a loss function that depends on the network weights and biases and finding a local or global minimum numerically. Some of the most popular algorithms are presented in the following sections. All presented algorithms produce a series of points in the space created by the network weights and biases. For simplicity, only weights are considered in this chapter. The point at the k-th optimization step is defined by a distinct set of weights and biases  $W^k$ , similar to coordinates in the state space.

### 3.4.1. Gradient Descent

In gradient descent, the next point  $W^{k+1}$ , i.e. the next set of weights and biases, is reached by computing the gradient of the loss function with respect to the network weights,  $\nabla_{W^k} L^k$ , and performing one step in the direction of steepest descent.  $W^{k+1}$  is calculated by the following equation:

$$W^{k+1} = W^k - \nabla_{W^k} L^k \cdot \alpha \quad (3.8)$$

with

$$\nabla_{W^k} L^k = \frac{1}{N} \cdot \sum_{n=1}^N \nabla_{W^k} L_n^k = \frac{1}{N} \cdot \sum_{n=1}^N [(a_n(s) - a_{n,desired}) \cdot \nabla_{W_k} a_n(s)] \quad (3.9)$$

and an arbitrary step size  $\alpha$ .

The choice of  $\alpha$  has a big effect on the convergence of the optimization algorithm. If  $\alpha$  is too small, each  $W_{k+1}$  is close to the respective  $W_k$  so it takes many steps to find a minimum. If  $\alpha$  is too big, the algorithm might not converge at all.

Ordinary gradient descent methods are not suitable for optimizing the output of large artificial neural networks because computing the gradient  $\frac{\partial L_k}{\partial W_k}$  for all the weights and biases is very computation-intensive and time-consuming.

### 3.4.2. Stochastic Gradient Descent

As mentioned before, the output of a neural network cannot be optimized efficiently with a gradient descent algorithm. One way to bypass this is Stochastic Gradient Descent. If the loss function is a sum like in eq. 3.7, it is sufficient to calculate the gradient of only one summand of the loss function, i.e. only consider one random training sample at each step  $k$  and perform a small step in the opposite direction:

$$\nabla_{W^k} L^k \approx \nabla_{W_k} L_n^k \quad (3.10)$$

$$W^{k+1} = W^k - \nabla_{W_k} L_n^k \cdot \alpha \quad (3.11)$$

If this is done for all samples successively, it also converges to the loss from eq.3.7. Thanks to the faster computation of  $\nabla_{W^k} L_n^k$ , SGD converges faster than ordinary gradient descent although eq.3.10 is only a rough estimate of the true gradient.

### 3.4.3. Momentum

Stochastic Gradient Descent has a fixed stepsize  $\alpha$ . Most advanced optimization algorithms have a variable stepsize and an update direction that depends on the previous updates to accelerate convergence. One example is to add a momentum term to the gradient from eq.3.11. This affects the step size and the direction of the update.

It takes into account the previous update and changes the direction of the next update. Like a ball rolling down a slope, the algorithm does not immediately change direction if the gradient of the slope changes direction or stop if the gradient is zero. If there is a plateau in the loss function or a weak local minimum, ordinary algorithms tend to get stuck there. With momentum, the risk of getting stuck is reduced.

### 3. Function Approximation

#### 3.4.4. The ADAM Algorithm

The ADAM-algorithm utilizes a modified momentum term for better convergence. Equation 3.12 shows one update step, the parameters  $\hat{m}_t$  and  $\hat{v}_t$  are calculated with equations 3.13 and 3.14.

$$W_{k+1} = W_k - \alpha \cdot \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon} \quad (3.12)$$

with

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} = \frac{\beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t}{1 - \beta_1^t} \quad (3.13)$$

and

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t} = \frac{\beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2}{1 - \beta_2^t} \quad (3.14)$$

The algorithm utilizes the first and second moments of the previous gradients to calculate the update. For the calculation of each  $m_t$  and  $v_t$ , the previous value  $m_{t-1}$  and the current gradient  $g_t$  are used.  $m_t$  and  $v_t$  are biased estimates of the first and second moments, there is a correction term in each of the denominators of equations 3.13 and 3.14 to account for that. See [4] for more details.

### 3.5. Overfitting

It is not always desirable to minimize the loss function to the absolute minimum of the loss function on the training data. Apart from the additional time it takes to get there, it increases the chances of running into problems.

When dealing with a continuous state space, it is common practice to discretize it and deal with the discretized state space instead. While a continuous state space consists of an infinite number of states, the number of states in the discretized state space is finite. This reduces computation time and makes some algorithms only applicable in the first place. If an appropriate grid is used (i.e. if it is dense enough in relevant areas of the state space), the solution is likely to be close to the solution of the continuous problem. The more grid points, the closer the solution is to the continuous solution.

If, however, a function approximator like a policy is trained on the discretized state set, it is possible that, although it might fit the training data very closely, it may not generalize well, i.e. give unexpected outputs when fed with states that have not been used in training. This problem is called *overfitting*. It especially occurs if the training data is noisy, i.e. each data point differs slightly from the expected value, and if a high number of parameters is trained with a limited dataset.

One way to deal with overfitting is splitting the state set into a set of training data and a set of validation data. After training on the training dataset, the value of the

### 3.5. Overfitting

loss function on the training data is compared to the value of the loss function on the validation data. If - from one training step to the next - the training loss decreases while the validation loss does not or even increases, overfitting is likely. This can be used as a criterion to stop the training.

The training set can also be split into three datasets, one for training, one for validation during training and the third one to check for overfitting after training. Typically, the third dataset contains approximately 20% of the samples, the remaining 80% are divided into training and validation data at the same ratio, so the parameter updates are done with 64% of the total sample count and 16% are used to check for overfitting during training.

Another way of preventing overfitting is to penalize large weights in the hidden layers.





## 4. Dynamic Programming




Figure 4.1.: classification of Machine Learning algorithms

This thesis deals with Dynamic Programming for policy optimization. Markus Zuern very successfully applied Trust Region Policy Optimization (TRPO) and Asynchronous Advantage Actor Critic (A3C) to a 2D trajectory optimization problem[6]. In his work, both algorithms deal with the complete environment, i.e. the glider dynamics and the stochastic wind distribution, at the same time. Whereas the wind distribution is unknown a priori, the glider dynamics are well known. Therefore it is possible to "pre-train" a policy in calm air and then use TRPO or A3C to improve its behavior when confronted with unknown factors such as thermal updrafts.

#### 4. Dynamic Programming




Figure 4.2.: Figure 3.1 von Markus

The term Dynamic Programming is used for optimization problems that can be decomposed into subproblems and solved by solving those subproblems. The solution to the original problem is then recomposed from the solutions of the subproblems. If the subproblems are very similar, the solution of one subproblem can be used to calculate another one. In such cases, dynamic programming is applicable.

In this case, finding the optimal action in each state is such a subproblem. If the algorithm is successful, the policy is optimal according to section 4.1. Unlike with TRPO and A3C, optimization is not done through sampling complete trajectories, but by iterating the state-values and policy for each state independently. After one sweep over the state space, the value function and policy are updated and the process is repeated.

In theory, starting TRPO or A3C with a policy that is already optimal in calm air should accelerate training with updrafts because the information about how to behave in a situation with no wind is already put into the policy.

## 4.1. The Principle of Optimality

According to the principle of optimality, the solution of some optimization problems can be put together from the solution of subproblems. This is the basis for all dynamic programming algorithms. If an agent chooses the optimal action in a state  $s_t$  and also in every successor state it visits  $s_{t+k}$  ( $k \in [1, T - t]$ ), the resulting trajectory is the optimal one from  $s_t$  to the terminal state.

The principle of optimality was first (todo: entdeckt/postuliert?) by Richard E. Bellman in 1954 (todo: belegen...)[1]. He proved that various optimization problems could be solved efficiently by making use of the Principle of Optimality instead of traditional optimal control theory. In his book *Applied Dynamic Programming*, Bellman covered numerous examples of optimal control problems and shows how they can be solved with Dynamic Programming [2].

(todo: das noch länger ausführen)

## 4.2. Types of Dynamic Programming Algorithms

### 4.2.1. Policy Evaluation

In Policy Evaluation, the goal is to find an approximation of the value function  $v_\pi$  that corresponds to a given policy  $\pi$ . The value of a state is defined as the expected total return from that state onwards until the episode terminates. Obviously, the expectation in equation 4.1 depends on the actions that are taken at each step and therefore on the policy that is followed. The state value that corresponds to a specific policy  $\pi$  is called  $V_\pi(s)$ .

$$V_\pi(s_t) = \mathbb{E}[G_t | s_t = s] = \mathbb{E} \left[ \sum_{k=0}^{T-t-1} \gamma^k r_{t+k+1} | s_t = s \right] \quad (4.1)$$

With a stochastic discrete policy, equation 4.1 becomes

$$V_\pi(s_t) = \sum_a \pi(a|s) (\mathcal{R}(s, a) + \gamma V_\pi(s')) \quad (4.2)$$

Equation 4.1 can be used to update the value of  $s_t$  with the expected return according to the current policy.

$$V_\pi(s_t) \leftarrow \mathbb{E}[G_t | s_t = s] = \mathbb{E} \left[ \sum_{k=0}^{T-t-1} \gamma^k r_{t+k+1} | s_t = s \right] \quad (4.3)$$

$$= \mathbb{E}[r_{t+1} + \gamma^1 r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots + \gamma^{T-t-1} r_T | s_t = s] \quad (4.4)$$

In 4.1, all terms except the first one can be replaced by the expected return from  $s_{t+1}$ :

## 4. Dynamic Programming

$$V_\pi(s_t) \leftarrow \mathbb{E}[r_{t+1} + \gamma G_{t+1} | s_t = s] \quad (4.5)$$

$$= \mathbb{E}[r_{t+1} + \gamma V_\pi(s_{t+1}) | s_t = s] \quad (4.6)$$

With a discrete stochastic policy, equation 4.6 becomes

$$V_{\pi, new}(s_t) \leftarrow \sum_a \pi(a|s)(r_t + \gamma V_\pi(s')) \quad (4.7)$$

In a state space with a finite set of states  $\mathcal{S}$ , writing down 4.7 for each state results in a system of  $|\mathcal{S}|$  linear equations that can be solved for  $V(s)$ . If the number of states is very large, this is not very efficient.

Instead, 4.1 is usually solved iteratively. The simplest way is to perform a one step lookahead from each state to get  $r_t$  and  $V(s_{t+1})$  from the current estimate of the state value function. Once all values are calculated, the values of all states are updated at the same time.

Each set of state values approximates the true value function of the problem better than the previous set of estimates. Every state value converges to the true state value under the given policy. It is not efficient to wait until the values have fully converged. Instead the process is repeated until the maximum absolute change in values lies beneath a certain threshold  $\epsilon$ .

### 4.2.2. Policy Iteration

Policy Iteration is an iterative way to calculate an estimate of the optimal policy of an MDP. Starting with an arbitrary policy and value function, one iteration of Policy Evaluation is performed to get an estimate of  $v_\pi$ . After that, a new policy is generated that chooses the greedy action with respect to  $v_\pi$  in each state. This alternating process of Policy Evaluation and Policy Improvement is repeated until the policy is satisfactory.

$$a_{greedy}(s_t) = \underset{a}{argmax} [Q(s_t, a)] \quad (4.8)$$

For a deterministic policy  $\pi$ , this yields:

$$\pi(a_{greedy}|s_t) = \mathbb{P}[a_t = a_{greedy}(s_t)|s_t] \leftarrow 1 \quad (4.9)$$

and for a stochastic policy  $\pi$ :

$$\mathbb{E}[\pi(s_t)] \leftarrow a_{greedy} \quad (4.10)$$

As mentioned in chapter 4.2.1, each policy evaluation step usually ends if the maximum difference between two subsequent values of the same state are sufficiently close.

There is a version of Policy Iteration where only one value update is done before updating the policy. This algorithm is called Optimistic Policy Iteration. Although the

first estimate of the value function is only a rough approximation, OPI also converges to the optimal value function and policy.

The performance of both algorithms on the glider-problem is compared in chapter 6.

### 4.2.3. Value Iteration

Value Iteration is similar to Policy Iteration. But instead of calculating a new policy at each iteration step, the state values  $V(s)$  are directly updated with the maximum possible successor value that is achievable from  $s$ . This yields a sequence of value functions not necessarily corresponding to a policy, but each one being a better estimate of  $v^*$  than its predecessor.

$$V(s_t) \leftarrow \max[r_{t+1} + \gamma V(s_{t+1}) | s_t = s] \quad (4.11)$$

This is equivalent to Optimistic Policy Iteration, where after one state value update at each state, the policy is replaced by the greedy policy with respect to the new value function. The only difference is that Value Iteration does not output an intermediate policy at each Iteration. Instead, the optimal policy can be calculated in the end by acting greedily with respect to the last value function.

## 4.3. The Contraction Mapping Theorem

A contraction mapping  $f : M \rightarrow M$  on a metric space  $M$  has the following property:

$$|f(x_1) - f(x_2)| \leq \gamma |x_1 - x_2| \quad (4.12)$$

with  $x_1, x_2 \in M$  and  $\gamma \in [0, 1)$ . Instead of  $|x_1 - x_2|$  and  $|f(x_1) - f(x_2)|$ , any measure of the distance between  $x_1$  and  $x_2$  or  $f(x_1)$  and  $f(x_2)$  can be used.

The term contraction comes from the fact that, graphically speaking, a contraction mapping reduces the distance between  $x_1$  and  $x_2$ .

### Discounted MDPs

Policy Evaluation in discounted MDPs is a contraction mapping. This result can be obtained in few steps. The mapping in equation 4.7 is known as the *Bellman Operator*  $T^\pi(v(s)) = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi v(s')$ . It maps from a state value function to a state value function. If the Bellman Operator is applied multiple times, this can be expressed by  $T^n(v)$ .

If the Bellman Operator is applied to two state value functions  $v_1(s)$  and  $v_2(s)$ , it reduces the distance between the two in value function space. This distance can be expressed by the supremum norm which is defined as follows:

$$\|v\|_\infty = \max_{s \in \mathcal{S}} |v(s)| \quad (4.13)$$

#### 4. Dynamic Programming

The supremum norm of  $v_1 - v_2$  is therefore

$$\|v_1(s) - v_2(s)\|_\infty = \|v_2(s) - v_1(s)\|_\infty = \max_{s \in \mathcal{S}} |v_1(s) - v_2(s)| = \max_{s \in \mathcal{S}} |v_2(s) - v_1(s)| \quad (4.14)$$

which is the biggest absolute difference of values of the same state in the state space. If the Bellman Operator is applied to both  $v_1$  and  $v_2$ , the contraction property proof is straightforward. For clarity, the arguments of  $v_1(s)$  and  $v_2(s)$  are omitted in the following equations.

$$\|T^\pi(v_1) - T^\pi(v_2)\|_\infty = \|(\mathcal{R}^\pi + \gamma \mathcal{P}^\pi v_1) - (\mathcal{R}^\pi + \gamma \mathcal{P}^\pi v_2)\|_\infty \quad (4.15)$$

$$= \|\gamma \mathcal{P}^\pi(v_1 - v_2)\|_\infty \quad (4.16)$$

$$\leq \|\gamma \mathcal{P}^\pi\| \|v_1 - v_2\|_\infty \quad (4.17)$$

$$\leq \|\gamma\| \|v_1 - v_2\|_\infty \quad (4.18)$$

$$= \gamma \|v_1 - v_2\|_\infty \quad (4.19)$$

Whenever  $s'$  happens to be the state where  $v_1$  and  $v_2$  differ the most, equality holds in line 4.17. If not, applying the supremum norm to  $v_1(s') - v_2(s')$  increases the value of the expression.  $\mathcal{P}^\pi$  must always be smaller than or equal to one. Similar to line 4.17, equality in line 4.18 holds only if  $\mathcal{P}^\pi = 1$ . Replacing  $\mathcal{P}^\pi$  by one in line 4.17 yields line 4.18. The left side of line 4.15 must thus be less than line 4.19 which proves the contraction property of the Bellman Operator in discounted MDPs.

#### Undiscounted MDPs

An MDP, where the discount factor  $\gamma$  is one, is called an undiscounted MDP. In such an MDP, the Bellman Operator looks like in equation 4.20.

$$T^\pi(v(s)) = \mathcal{R}^\pi + \mathcal{P}^\pi v(s') \quad (4.20)$$

If  $\gamma$  is one, the above proof of the contraction property does not hold anymore. It can however be proven, that the Bellman Operator still contracts. In the undiscounted case, this takes more than one step. Equation 4.21 shows the n-step contraction property. For clarity, the index  $\pi$  and the argument  $s$  are omitted.

$$\|T^{(n)}(v_1) - T^{(n)}(v_2)\|_\infty \leq \|v_1 - v_2\|_\infty \quad (4.21)$$

Equation 4.21 means that applying the Bellman Operator  $n$  times to two value functions brings them closer together in value function space. As a result, the Bellman Operator also leads to  $v_1 \approx v_2$  eventually in undiscounted MDPs.

The following proof is more sophisticated than for discounted MDPs. (todo: Proof)

## 5. Trajectory Optimization with Policy Iteration

The goal of this work is to find an optimal trajectory of a glider moving through calm air and reaching a given distance in minimal time. The resulting optimal policy can be used as a starting point for training the glider on scenarios with different wind conditions utilizing other algorithms like reinforcement learning. The basic idea is to use the known information about the glider dynamics to pre-train a policy that is - in the next step - improved by reinforcement learning.

### 5.1. Glider Representation

In the context of this work, the aircraft is treated as a mass point. Rotational dynamics are neglected. This is sufficient for trajectory optimization algorithms. See [3] for more details.

### 5.2. 2D Environment

In this scenario, the glider moves in the geodetic vertical plane. Its control is the angle of attack  $\alpha$ . When  $\alpha$  is increased, the lift  $L$  increases. This results in an increase of  $\dot{\gamma}$ , as can be seen in equation 5.4. This way the agent can control its velocity vector and therefore its position. All other parameters of the glider are shown in appendix A. Any 3D-Scenario where the horizontal trajectory is restricted by obstacles or the optimal ground trace is obvious can be regarded as a 2D problem. The only control variable to be optimized is then the angle of attack or an equivalent quantity (vertical acceleration, flight path angle,...).

#### 5.2.1. Policy Representation

As the Policy Iteration algorithm is running on a rectangular grid and the states that are updated do not change, a table that stores the greedy action for every state is sufficient for the Policy Iteration algorithm. Once the optimal policy is found, a neural network is trained to approximate the optimal action for each of the states on the grid. This policy is then used to do TRPO.

### 5.2.2. Equations of Motion

The physical state of the glider in the geodetic vertical plane consists of the position  ${}_g\mathbf{r} = [x, z]^T$  and velocity  ${}_g\mathbf{v} = [{}_g u_K, {}_g w_K]^T$ . Apart from the geodetic reference frame (index g), a body fixed frame (index f), a trajectory fixed frame (index k) and an aerodynamic frame (index a) are used to describe the glider dynamics. In the vertical plane, coordinates can be transformed between the frames by a translation and a rotation around the y-axis. Figure 5.1 illustrates the transformation angles that connect the reference frames.




Figure 5.1.: Angles in three-dimensional flight.

In the vertical plane, there are 4 state variables,  $x$ ,  $z$ ,  $u = \dot{x}$  and  $w = -\dot{z}$ , each with respect to a geodetic reference frame that has its origin at a point on the earth surface. The aerodynamic forces are often expressed in an aerodynamic reference frame. The coordinate transformations  $u$  and  $w$  can also be expressed by  $V$  and  $\gamma$ .

The dynamics in the vertical plane are given by the following equations. They are taken from [3].



$$\dot{x} = V \cdot \cos\gamma \quad (5.1)$$

$$\dot{z} = -V \cdot \sin\gamma \quad (5.2)$$

$$\dot{V} = -\frac{D + g \cdot \cos\gamma}{m} \quad (5.3)$$

$$\dot{\gamma} = \frac{L}{V \cdot m} - \frac{g \cdot \cos\gamma}{V} \quad (5.4)$$

where  $L = q \cdot c_L \cdot S$ ,  $D = q \cdot c_D \cdot S = \frac{\rho}{2} \cdot V^2 \cdot c_d \cdot S$  and  $V = \sqrt{u^2 + w^2}$ .

### 5.2.3. Maximum Range in a given Configuration

An airplane with no propulsion system moving through calm air reaches its maximum range when flying at the maximum ratio of lift to drag. Lift and drag can be expressed as the product of dynamic pressure  $q$ , a reference area  $S$  and a dimensionless coefficient  $c_L$  or  $c_D$ :

$$L = q \cdot S \cdot c_L = \frac{\rho}{2} v^2 \cdot S \cdot c_L \quad (5.5)$$

$$D = q \cdot S \cdot c_D = \frac{\rho}{2} v^2 \cdot S \cdot c_D \quad (5.6)$$

$$\frac{L}{D} = \frac{c_L}{c_D} \quad (5.7)$$

The optimal lift to drag ratio is also the optimal ratio of  $c_L$  and  $c_D$ :

$$\left(\frac{L}{D}\right)_{max} = \left(\frac{c_L}{c_D}\right)_{max} = \left(\frac{c_L}{c_D}\right)_{opt} \quad (5.8)$$

With a symmetric drag polar model, this yields:

$$c_{L,opt} = c_L \left( \left( \frac{c_L}{c_D} \right)_{opt} \right) = \sqrt{\frac{c_{d0}}{k}} \quad (5.9)$$

and

$$\alpha_{opt} \approx \frac{\Lambda + 2}{2\pi\Lambda} \cdot c_{L,opt} \quad (5.10)$$

A policy that is trained to return  $\alpha_{opt}$  for every state can be used as a starting point for policy optimization with any policy optimization algorithm.

## 5. Trajectory Optimization with Policy Iteration

### 5.2.4. Following the Line of Sight Towards the Target

At each state in the state space, the line of sight from the agent to the target can be expressed by a unit vector pointing from the agent towards the target:

$$\vec{e}_{los} = \frac{\vec{r}_{target} - \vec{r}_{agent}}{|\vec{r}_{target} - \vec{r}_{agent}|} \quad (5.11)$$

The angle  $\gamma_{los} \in [-\pi, \pi]$  between the horizontal axis and  $\vec{e}_{los}$  can be expressed by the coordinates of the target and the agent.

$$\gamma_{los} = \text{atan2}\left(\frac{\Delta z}{\Delta x}\right) = \text{atan2}\left(\frac{z_{target} - z_{agent}}{x_{target} - x_{agent}}\right) \quad (5.12)$$

$\gamma_{los}$  is positive if the agent is below the target.

The velocity vector  $\vec{v} = [u, w]^T$  represents the direction in which the glider is moving. Like  $\gamma_{los}$ , the flight path angle can be calculated with the atan2-function.

$$\gamma = \text{atan2}\left(\frac{u}{w}\right) \quad (5.13)$$

### 5.2.5. Discretization of the State- and Action Space

The state space is discretized with a rectangular grid, i.e. it is replaced by a set of points  $s_n = [x_i, z_j, u_k, w_l]$  evenly distributed across all dimensions.

In the 2D-scenario, the action space is one-dimensional, the only action is the angle of attack  $\alpha$ . For the policy improvement step, a finite number of evenly spaced actions is sampled from the infinite set of possible actions between  $\alpha_{min} = 0$  and  $\alpha_{max} = 0.2$ . Assuming that the mapping from actions to returns is continuous, the action that yields the maximum reward (plus successor state value) out of the sampled actions is an approximation of the true greedy action.

By discretization, the continuous trajectory optimization problem is made time- and space-discrete. For the purpose of keeping calculation time low, a sampling rate  $\Delta t$  of 1s is used. With a speed of  $15 - 25 \frac{m}{s}$ , the agent covers about  $15 - 25m$  within one time step. The grid points should be separated by approximately that distance in order to make tabular solution methods feasible. For the different scenarios, this yields the grid resolutions shown in table 5.1.

## 5.3. 3D Environment

todo

### 5.3.1. Discretization of the state and action space

todo




Figure 5.2.: The discretized state space. For one set of  $x_3$  and  $z_4$ , the grid for the speed vector is drawn.

distance	500m	1000m	2000m
$n_x$	25	50	100
$n_z$	16	16	16
$n_u$	7	7	7
$n_w$	7	7	7

Table 5.1.: Grid parameters for trajectory optimization



## 6. Results

### 6.1. 2D Policy Iteration

In the 2D scenario, three Dynamic Programming algorithms are compared, Generalized Policy Iteration with up to 1000 evaluation iterations before performing the Policy Improvement step, Optimistic Policy Iteration and Value Iteration. As OPI and VI are technically equivalent, the interesting part here is which algorithm yields a nearly optimal policy in less time. In the GPI method, the iterative evaluation stops if the values of all states change by less than one from one evaluation step to the next.




Figure 6.1.: Trajectory and angle of attack after 100 Optimistic Policy Iterations

## 6. Results




Figure 6.2.: Trajectory and angle of attack after 100 Generalized Policy Iterations with up to 100 evaluation steps in each iteration

Figure 6.1 shows the trajectory in scenario 1 after 100 iterations of Optimistic Policy Iteration. The agent is dropped at the start state with zero velocity. As can be seen, both policies direct the agent towards the target.

Figure 6.4 shows the flight times from the start-state to the goal after each Iteration of GPI, OPI and VI.

Although all algorithms theoretically take up to  $|\mathcal{A}||\mathcal{S}|$  iterations to converge (Beleg!), they are known to typically converge after surprisingly few iterations.

During each Policy Iteration step with GPI, it takes between 60 and 70 Policy Evaluation Steps until approximate convergence. Therefore each Policy Iteration takes significantly longer in the GPI case than with Optimistic Policy Iteration where Evaluation is stopped after only one Evaluation Iteration.

Both algorithms yield policies that are able to find the goal from various initial states.




Figure 6.3.: Trajectory and angle of attack after 100 Value Iterations

## 6.2. 3D Scenarios

## 6. Results




Figure 6.4.: Trajectory and angle of attack after 100 iterations of Generalized Policy Iteration with 100 evaluation steps in each iteration



## 7. Discussion

Platzhalter...



## A. Appendix A: Glider Parameters and Scenario Data

mass $m$	$3.366kg$
wing area $S$	$0.568m^2$
aspect ratio $\Lambda$	$10.2$
oswald efficiency factor $e$	$0.9$
$c_{D0}$	$0.015$

Table A.1.: glider data

## A. Appendix A: Glider Parameters and Scenario Data

scenario	1	2	3
distance	500m	1000m	2000m
start state	$s_0 = [0, 100, 0, 0]^T$	Platzhalter	Platzhalter

Table A.2.: scenario data

## B. Appendix B: Computer Configuration

Table B.1 contains the technical data of the system used for calculations.

CPU	Intel i7-7700HQ @ 2.80 GHz (Turbo Boost up to 3.80 GHz, 4 physical cores, hyper threading)
RAM	16 GB DDR3 RAM @ 3200 MHz
GPU	Nvidia Geforce GTX1050Ti (4GB VRAM, 768 CUDA cores)
OS	Ubuntu 16.04 LTS

Table B.1.: computer specs



# C. Appendix C: Standardization Values

## C.1. 2D Scenarios

distance	$\mu$	$\sigma$
500m	Platzhalter	Platzhalter
2000m	Platzhalter	Platzhalter

Table C.1.: Standardization values for 2D trajectory optimization





# Bibliography

- [1] BELLMAN, R. E.: The Theory of Dynamic Programming. (1954)
- [2] BELLMAN, R. E.: Applied Dynamic Programming. (1962)
- [3] FICHTER, W. ; GRIMM, W. : *Flugmechanik*. Shaker Verlag, 2009
- [4] KINGMA, D. P. ; BA, J. : Adam: A Method for Stochastic Optimization. In: *CoRR* abs/1412.6980 (2014). <http://arxiv.org/abs/1412.6980>
- [5] SUTTON, R. S. ; ; BARTO, A. G.: *Reinforcement Learning - An Introduction*. The MIT Press, 2018
- [6] ZUERN, M. : *Autonomous Soaring through Unknown Atmosphere, Master Thesis*. Institute for Flight Mechanics and Control, Stuttgart University, 2017